

PRIVATE INTERVIEW PREPARATION

CensusChat

Engineering Interview Manual

How to explain the system, defend the tradeoffs,
and map the build to the Snowflake Applied AI rubric.

01

LEAD WITH

Deterministic SQL trust boundary

02

TEACH

Closed topology, open vocabulary

03

BE CANDID

Grounding evidence is not a runtime guarantee

SECTION 01

The 90-second opening

Start with the problem, the governing architecture, and one honest boundary.

SAY THIS CensusChat is a production-minded chat agent for the 2020 ACS five-year estimates. It combines runtime discovery over thousands of Census variables with a small, explicit tool loop. The model can decide what to look up, but it cannot decide what SQL is safe: every request-time Snowflake query crosses a default-deny AST gate. The design optimizes for grounded answers, honest ambiguity, and observable failure under a 24-hour constraint.

The story in three moves

- **Coverage without prompt bloat.** The schema topology is small and stable; the variable vocabulary is large and discovered at runtime through local FTS.
- **Autonomy inside deterministic boundaries.** Sonnet selects among exactly three tools, while code enforces SQL safety, unresolved geography ambiguity, retry limits, and terminal stream behavior.
- **Evidence with named limits.** Offline tests prove deterministic layers. Live evals exercise the real model and Snowflake stack. Semantic prose quality remains human-reviewed.

censuschat

Ask about US population, income, or housing from the 2020 ACS 5-year estimates.

Chat

How It Works

Evidence

Evals

Example questions

New Chat

Start with a factual question

Population of Harris County, Texas?

Follow up after the first question

What about households?

See ambiguity handling

How many people live in Washington County?

See an unsupported request

How many people will live in Texas in 2050?

e.g. What's the population of Travis County, Texas?

Send

Current reviewer surface: Chat, How It Works, Evidence, Evals, plus a New Chat boundary for session isolation.

DO NOT OVERCLAIM The strongest hard guarantee is SQL safety. Numeric answer grounding is a model instruction with selected offline evidence checks, not a serving-time validator for every final-answer number.

App snapshot 606e35b | Verified 2026-08-12

Do not overstate soft guarantees

02

SECTION 02

Assignment rubric map

Tie each rubric dimension to implementation evidence, then volunteer the relevant limitation.

Dimension	Evidence	Known limit	Interview move
LLM / AI Engineering	Runtime variable and geography discovery; three-tool Sonnet loop; schema rules in context; deterministic ambiguity backstop; AST SQL gate.	Final-answer numbers are not all runtime-validated. Prose quality has no calibrated model judge.	Lead with the split between model reasoning and code enforcement.
Production Quality	SSE streaming; graceful degraded mode; sanitized errors; bounded recovery; SQL statement timeout; durable Evidence traces; health endpoint.	Single-instance SQLite, no rate limit or spend cap, cached Snowflake health, soft watchdog.	Explain failure behavior before listing infrastructure.
Judgment Under Constraints	2020-only allowlist; FTS instead of embeddings; no agent framework; three tools only; decennial and Langfuse cut deliberately.	The interface arrived late in the original build. Some broader capability remains flaky or informational.	Defend what was excluded as strongly as what was built.
Reflection and Self-Awareness	The reflection names false-green tests, prompt defects, eval-calibration failures, production gaps, and a prioritized next-day plan.	Do not turn reflection into self-criticism. Use incidents to show changed engineering judgment.	Tell one concrete failure and the general lesson it produced.

WHY IT MATTERS

The assignment explicitly evaluates the ability to explain and defend open-ended technical decisions. A rubric answer should therefore have the form: choice, evidence, tradeoff, rejected alternative.

A compact evidence ladder


- **Code:** the strongest evidence for deterministic guarantees.
- **Offline tests:** repeatable proof that code-level contracts hold.
- **Live evals:** evidence about the integrated model, prompt, tools, and database.
- **Reflection:** evidence of judgment, not evidence that a behavior works.

SOURCE PATHS docs/assignment.pdf pp. 2-5 | README.md | docs/reflection.md

SECTION 03

Architecture at a glance

The model is an orchestrator inside a narrow action space, not the security boundary.



The governing design insight

CensusChat separates **closed topology** from **open vocabulary**. Join rules, Census block-group key structure, rollup recipes, and statistical constraints are small enough to teach in the system prompt. Thousands of variable labels and geography records stay outside the prompt and are retrieved only when needed.

SAY THIS

The prompt teaches how the database works, not everything the database contains. That keeps context small while preserving broad variable coverage.

Why three tools is a feature


- **search_census_variables** translates user language into Census field identifiers using local SQLite FTS5.
- **resolve_geography** maps a state or county name to deterministic FIPS candidates using local SQLite.
- **run_census_sql** is the only request-time Snowflake code path, and every call crosses the SQL gate.

SOURCE PATHS `src/agent.py` | `src/tools.py` | `src/sqlgate.py` | `src/model_config.py`

SECTION 04

One request, end to end

Walk through the turn in time order. This is the clearest way to explain state, tools, safety, and observability together.



What happens at each boundary

- The browser sends a client-generated **session_id** and message to FastAPI, which returns an SSE stream.
- The agent replays the full user/assistant history from SQLite, then gives the classifier only the two most recent stored messages for context-aware routing.
- Sonnet discovers variables and geography locally, proposes SQL, receives normalized query rows, and streams answer tokens.
- Each tool call emits start/end events. Every turn terminates with **done** or **error**. A separate trace record stores spans, final answer, terminal status, and timing.

WHY IT MATTERS	The data flow makes the trust boundaries inspectable: user text becomes model context, never SQL string interpolation; proposed SQL becomes executable only after structural validation.
DO NOT OVERCLAIM	Full-history replay is simple and effective for a demo, but it can replay orphaned user turns after infrastructure failures. New Chat gives the user a fresh session boundary; a production fix would make turn persistence transactional or exclude failed turns from replay.

SOURCE PATHS `src/app.py` | `src/agent.py` | `src/sessions.py` | `src/tracing.py`

SECTION 05

Data model and statistical correctness

Many plausible-looking Census answers are wrong because the aggregation is invalid, not because the SQL failed.

Concept	Implementation
Source	SafeGraph Open Census Data share on Snowflake Marketplace, 2020 ACS 5-year estimates covering 2016-2020.
Physical grain	Census block group. State and county answers roll up CBG rows using FIPS prefixes.
Supported geography	State and county. No city/place, ZIP, or metro boundary exists in the shipped source.
Searchable vocabulary	All supported estimate fields are indexed locally; margin-of-error and B99 allocation rows are excluded from retrieval to avoid plausible near-miss answers.
Vintage	2020 only, enforced by the SQL table allowlist rather than by a prompt preference.

The aggregation rules to explain

- **Counts roll up by SUM.** State and county filters use prefixes of the 12-character CENSUS_BLOCK_GROUP code.
- **Medians never aggregate.** SUM or AVG of block-group medians is statistically invalid. The agent may offer a true mean only when a valid aggregate numerator and denominator exist.
- **Universes must match.** Population, households, workers, and occupied units are not interchangeable denominators.
- **NULL is not zero.** For direct or simply aliased variables with known lineage, NULL becomes 'not reported' and the verified income top-code becomes '\$250,000 or more.' Derived and aggregate projections remain unchanged at that seam.

SAY THIS

The system treats statistical validity as data and code wherever possible. A query can be syntactically safe and still be analytically wrong, so retrieval metadata tells the agent which variables can roll up.

Why 2015-2019 was not added

The two available ACS five-year releases overlap in four of five years, and 2020 redrew block-group boundaries. Adding the older tables would double retrieval ambiguity without making a clean trend comparison valid. If multi-vintage access became a requirement, the next design would add vintage to retrieval results and reject cross-vintage statements in code.

SOURCE PATHS docs/schema-notes.md | docs/decisions.md D-003/D-005/D-008 | src/tools.py

SECTION 06

Trust boundaries and graceful failure

Separate advisory model behavior from code-enforced invariants. This is the strongest system-design section.

Layer	Kind	Role	Failure behavior
System prompt	Soft	Quoting, aggregation, grounding, and answer-style instructions.	The model may ignore or misapply it.
Haiku classifier	Soft	Fast-fail off-topic, adversarial, and inappropriate input.	Fails open on timeout or error.
Ambiguity backstop	Hard	Blocks SQL when a geography resolved as ambiguous during the turn.	Produces deterministic clarification.
SQL gate	Hard	Parses one SELECT, checks every table, rejects banned structures, injects LIMIT.	Defaults to deny; rejected SQL never executes.
Snowflake session	Hard	25-second statement timeout and sanitized SQL only.	Bounds database execution, not connection or model time.

What the SQL gate checks

- sqllglot parses with the Snowflake dialect; regex is not the parser.
- Exactly one statement, SELECT only, with CTEs allowed and DDL, DML, INTO, calls, session variables, and unverifiable table positions rejected.
- Every referenced table must match the fully qualified allowlist. Explicit columns are required; COUNT(*) is the narrow exemption.
- A LIMIT is injected when missing. User text is never interpolated into SQL.

SAY THIS

The classifier is allowed to fail open because it is not load-bearing. Availability stays high, while the SQL boundary still fails closed. At scale I would add defense in depth with a least-privilege Snowflake role and resource monitor.

Bounds that prevent hanging or uncontrolled repair

- Two retries after SQL failure or zero rows; then deterministic honest failure.
- Eight total tool-loop rounds regardless of success or failure.
- A 50-second watchdog stops new rounds after the budget, but cannot cancel a call already in flight.

SOURCE PATHS `src/sqlgate.py` | `src/guardrail.py` | `src/agent.py` | `src/contracts.py`

SECTION 07

State, streaming, and Evidence

Explain these as three separate responsibilities, even though the UI brings them together.

Responsibility	Mechanism	Why
Conversation state	SQLite messages keyed by session_id	Full user/assistant replay gives multi-turn continuity and survives reloads.
Transport state	SSE ChatEvents	Tokens, tool start/end, terminal done/error, and progress reach the browser incrementally.
Evidence state	Separate SQLite trace store	Guardrail, model, and tool spans plus final answer, terminal status, timing, and raw JSON.
Browser boundary	localStorage session_id + New Chat	New Chat starts a fresh replay context without deleting old Evidence history.

What to point out in the Evidence tab

- The guardrail verdict and latency show whether the turn took the fast-fail path.
- Model spans expose stop reason, token counts, and time spent per round.
- Tool spans show arguments, bounded result summaries, success/failure, and latency.
- The terminal step pairs the final answer with done/error status and total timing.

WHY IT MATTERS

Observability is part of the reviewer story because it makes a stochastic system inspectable. It also revealed the recent session-replay bug: traces showed an older Washington County turn being replayed into a Harris County request.

DO NOT OVERCLAIM

Evidence is application-local tracing, not Langfuse. It has no cross-service search, alerting, retention policy, or replica-aware aggregation. Recording is deliberately fail-soft so a tracing bug cannot break chat.

SOURCE PATHS [src/sessions.py](#) | [src/tracing.py](#) | [static/index.html](#) | [docs/decisions.md](#) D-021/D-023/D-025

SECTION 08

Testing and evaluation strategy

Use the right instrument for each layer. Do not ask mocked unit tests to certify model behavior.



Suite	Scenarios	Contract
Regression (6)	DF-05, MT-01, AMB-01, UN-01, OT-01, INJ-02	Stable, objectively scored behaviors. Two live trials must both pass when the manual workflow is run.
Capability (8)	DF-01, CMP-01, AMB-02, PM-02, PM-03, AMB-03, UN-08, PM-08	Broader or stochastic coverage. Pass, fail, or inconclusive, but non-blocking.

What deterministic graders can prove

- Expected variable and geography appear in captured tool evidence.
- Answer is nonblank or contains a stable reference value.
- Refusal paths avoid tools and use refusal language; ambiguity paths ask and avoid SQL.
- Protected median variables are not aggregated with SUM or AVG.
- Four-digit answer figures are checked only against visible final-turn query-row cells; hidden rows become inconclusive rather than passing.

DO NOT OVERCLAIM

The committed 14/14 artifact is from commit d44c1cc on 2026-08-06 and predates the current suite and tri-state contracts. The UI labels it legacy. It demonstrates that a real-stack run existed, not that the current branch is 14/14.

SAY THIS

I did not add an LLM judge because an uncalibrated judge would create a score, not evidence. Prose quality stays human-reviewed until a judge is compared with labeled examples and agreement is measured on held-out data.

SOURCE PATHS tests/ | evals/scenarios.py | evals/run_evals.py | .github/workflows/

SECTION 09

Deployment and production quality

The app is deployable and observable as a controlled demo. Production at customer scale would require a different state and cost-control layer.

Layer	Choice	Operational meaning
Edge	Native Caddy on EC2	TLS, basic auth, SSE flushing; app port bound to loopback on the host.
Application	Dockerized FastAPI/Uvicorn	One stateless process around stateful local files; image includes frontend and committed eval artifacts.
Model	Anthropic SDK	Sonnet agent and Haiku classifier, pinned in one module.
Data	Snowflake Marketplace share	Only run_census_sql touches it at request time; local snapshot handles discovery.
Persistence	Mounted SQLite files	Sessions and traces survive restart and deploy, but do not support horizontal replicas.
CI	GitHub Actions	Credential-free offline tests on PR/main; protected, manually triggered paid regression workflow.

Verified deployment snapshot

EVIDENCE

https://censuschat.brianmar.com and the local instance both returned status=ok, snapshot=ok, snowflake=ok on 2026-08-12. This is a point-in-time check, not continuous availability evidence.

What is still missing for production

- Rate limiting, per-session spend caps, abuse monitoring, and a least-privilege Snowflake role with a resource monitor.
- Shared session/trace storage for multiple replicas, plus retention and deletion policies.
- Per-call cancellation deadlines. The current watchdog checks only between rounds.
- Live health or circuit breaking that can detect Snowflake failure after boot without violating the one-query-path invariant.
- Central telemetry with search, alerting, and deployment correlation.

SAY THIS

I would ship this as a time-bounded reviewer demo. I would not call it multi-tenant production infrastructure until cost controls, shared state, cancellation, and centralized observability are added.

SOURCE PATHS Dockerfile | docker-compose.yml | deploy.sh | src/health.py | .github/workflows/

SECTION 10

Five-minute reviewer walkthrough

Use one coherent session. Show behavior first, then open one evidence surface to explain why it happened.

Step	Prompt/action	What it demonstrates	What to inspect
1	Population of Harris County, Texas?	Variable search, exact county resolution, safe state/county rollup, grounded fact.	Point to tool order and final QueryResult evidence.
2	What about households?	Full-history replay carries Harris County into the follow-up without restatement.	Show the second variable resolution and same geography.
3	How many people live in Washington County?	Thirty matches should trigger clarification, not a silent choice.	No Snowflake SQL while ambiguity is unresolved.
4	How many people will live in Texas in 2050?	The ACS is historical measurement, not projection. Fast refusal should avoid SQL.	Guardrail or answer path, zero database work.
5	Open Evidence, then Evals	Evidence explains one real turn. Evals explains test intent and result provenance.	Call out legacy benchmark labeling before the reviewer does.

Suggested narration

- Before the first prompt: 'I will show a fact, a context-dependent follow-up, ambiguity, and an unsupported request.'
- Before Evidence: 'The chat is the product surface; Evidence is the explanation surface.'
- Before Evals: 'The app does not pretend every behavior has the same measurement quality, so stable regression and broader capability are separate.'

DO NOT OVERCLAIM

A live model can vary. If a scenario behaves unexpectedly, use Evidence to localize the failure and discuss the appropriate fix. Do not rerun until it happens to look good.

Recovery line if the demo fails

SAY THIS

This is exactly why the trace and eval layers exist. Let me show whether the failure was routing, discovery, SQL validation, execution, or answer generation, and then I will explain which layer should own the fix.

SOURCE PATHS README.md: Evaluating the running demo | static/index.html | evals/README.md

SECTION 11

Judgment under the 24-hour constraint

The strongest answer is not 'I built everything.' It is 'I knew which guarantees were worth encoding and which breadth to cut.'

Decision	Why it was rational	Alternative rejected
FTS5 over embeddings	Variable retrieval needed exact, local, inspectable matching. A probe showed a morphology/ranking problem, not missing semantic representation.	Embedding dependency, index build, retrieval opacity, and tuning time.
Handwritten tool loop	Exactly three tools and explicit SSE/tracing needs kept control flow small and legible.	Framework abstraction and repair behavior that would be harder to inspect.
2020 only	Older ACS release overlaps four years and uses changed CBG boundaries. Code-enforced scope prevents invalid comparisons.	Broader but misleading vintage coverage.
SQL gate before feature breadth	A single unsafe or structurally unverifiable query is a customer risk. Default-deny validation is deterministic and testable.	Prompt-only SQL safety or a regex denylist.
SQLite sessions and traces	Fast durable evidence for one demo instance with no external service dependency.	Distributed state infrastructure that the assignment did not require.
No LLM judge	Semantic scores are not trustworthy until calibrated against human labels.	A polished but unvalidated quality number.

Two incidents that changed the engineering approach

- **Prompt defect hidden by recovery.** Unquoted mixed-case Census identifiers failed live even while mocked tests passed. Successful self-repair initially made the systematic defect look like resilience. Lesson: recovery attempts are failure signals and real-stack evals test the product, not just the code.
- **Uncalibrated grounding check produced false reds.** Variable IDs and missing division logic were misread as fabrication. Lesson: an eval check is a measuring instrument and needs known-good calibration plus mutation tests.

SAY THIS

My best judgment was encoding high-cost mistakes as deterministic boundaries: unsafe SQL, unresolved geography, invalid median aggregation, bounded repair, and honest terminal behavior. I cut breadth where it would weaken those guarantees.

SOURCE PATHS docs/reflection.md | docs/decisions.md | docs/01-architecture.md

SECTION 12

Likely interview questions, part I

Keep each answer under a minute. Name the tradeoff and the rejected alternative.

Why does the classifier fail open?

Because it is an availability and UX layer, not the trust boundary. If Haiku is unavailable, a legitimate Census request should still proceed. SQL safety remains fail-closed below it, and off-topic content still cannot escape the three-tool action space. The tradeoff is that an outage may allow irrelevant turns to reach Sonnet and cost more.

Why FTS5 instead of embeddings?

The retrieval target is a structured vocabulary of Census labels and IDs, and the observed failures were token morphology and ranking. FTS5 is local, deterministic, cheap, and inspectable. I would add embeddings only after an error analysis showed semantic recall failures that lexical retrieval could not solve.

Why did you avoid LangChain or LangGraph?

The loop has three tools, one model vendor, eight bounded rounds, and custom streaming and trace events. A handwritten Anthropic loop makes every transition and failure rule visible. A framework would become attractive when orchestration complexity, provider portability, or durable workflow resumption exceeded this small control surface.

Why only the 2020 ACS five-year vintage?

The older release overlaps four of five years and uses different block-group boundaries. Adding it would create plausible but invalid trend comparisons. The current allowlist makes cross-vintage SQL impossible. If multi-vintage access became necessary, I would add vintage-aware retrieval and a gate rule rejecting mixed-vintage statements.

Why can counts aggregate but medians cannot?

Counts are additive across non-overlapping block groups. A median is an order statistic and cannot be reconstructed from subgroup medians alone. Where aggregate numerator and denominator variables exist, the system may compute a true mean and explicitly state that substitution.

Why exactly three tools?

Each tool owns one distinct external effect: discover variables, resolve geography, execute data SQL. That is enough for the assignment while keeping the action space auditable. Adding a tool should require a new responsibility that cannot fit an existing boundary, not merely convenience.

Why is the 50-second watchdog soft?

It is checked between agent rounds, so it can stop new tool work after the budget is observed, but it cannot interrupt an Anthropic or Snowflake call already in flight. That leaves headroom under the assignment's 60-second target without creating a guarantee. A production version would add per-call cancellation and a request-wide deadline.

**INTERVIEW
PATTERN**

Answer architecture questions as: requirement, choice, why it fits this system, tradeoff, threshold for revisiting the choice.

SECTION 13

Likely interview questions, part II

These answers distinguish demonstrated behavior from unfinished production work.

How do you prove that answers are grounded?

Today I can prove selected parts: expected tools ran, expected IDs were resolved, and visible answer figures match captured final-turn query-row cells in the eval harness. I cannot claim every final-answer number is runtime-validated. The strongest next step is to retain complete per-turn query evidence, buffer the final answer, and validate structured numeric claims before release.

Why have an inconclusive eval outcome?

It prevents missing evidence from becoming a false pass. Bounded trace summaries may hide later query rows, so some grounding checks cannot honestly decide. Inconclusive stays non-passing and in the denominator. It is informational in capability reporting and blocking in regression trials.

Why fourteen scenarios, and why not more?

The current bottleneck is grader quality, not case count. Six stable regression scenarios protect objectively measurable behavior; eight capability scenarios expose broader or stochastic behavior without creating merge roulette. I would expand only with cases tied to observed failures and checks that have been proven to fail correctly.

What breaks first under real traffic?

Cost exposure and state topology. There is no rate limit or spend cap, and SQLite sessions and traces are local to one instance. Next come cancellation, stale boot-time health, and observability. I would add limits and shared state before adding replicas because replication without shared state would make behavior less coherent.

How did you use AI coding tools?

AI produced much of the implementation and independent review agents found real defects. My responsibility was dataset recon, architecture, constraints, acceptance criteria, live verification, and deciding which suggestions to reject. The key lesson was that code review agents still missed a prompt-plus-database integration defect, so I added real-stack eval evidence rather than assuming more review would solve it.

What would you build next?

First, serving-time numeric provenance. Second, transactional session turns so failed requests cannot remain as orphaned instructions. Third, per-call cancellation. Fourth, rate limits, spend caps, shared state, and centralized telemetry. Fifth, a genuinely independent decennial source for conflicting-answer cases.

DO NOT
OVERCLAIM

Do not say 'production ready' without a qualifier. Say 'production-minded reviewer demo' and immediately name the scale, cost, and observability work required for customer traffic.

SECTION 14

Interview cheat sheet

Read this page immediately before the interview. It is the compressed version of the entire manual.

FIVE CLAIMS TO LEAD WITH

- The AST SQL gate is the hard trust boundary.
- Closed topology plus open runtime vocabulary avoids schema prompt bloat.
- Exactly three tools keep the action space narrow and auditable.
- Ambiguity, retries, and terminal behavior have code backstops.
- Tests, live evals, and human review measure different layers.

KEY NUMBERS

- 2020 ACS five-year estimates, covering 2016-2020.
- 3 tools: variable search, geography resolution, SQL execution.
- 2 recovery retries, 8 total rounds, 50-second soft watchdog.
- 25-second Snowflake statement timeout, 200-row injected limit.
- 6 regression + 8 capability scenarios.
- 449 offline tests at the documented snapshot.

FIVE CLAIMS NOT TO MAKE

- 'Every answer number is runtime-grounded.'
- 'The 14/14 artifact proves the current branch.'
- 'The watchdog guarantees the 60-second requirement.'
- 'Evidence is production observability.'
- 'The app is horizontally scalable or cost bounded.'

NEXT FIVE IMPROVEMENTS

- Validate structured numeric claims before final answer release.
- Make turn persistence transactional; exclude failed orphan turns.
- Add per-call deadlines and cancellation.
- Add rate limits, spend caps, shared state, and central telemetry.
- Add an independent decennial source for conflict handling.

File map for live discussion

Question	Open
How does the loop work?	src/agent.py
Where is the hard boundary?	src/sqlgate.py
How are variables and geographies found?	src/tools.py + src/snapshot.py
How is context stored?	src/sessions.py
How is a turn explained?	src/tracing.py + Evidence tab
What is actually evaluated?	evals/scenarios.py + evals/README.md
What changed and why?	docs/decisions.md + docs/reflection.md

LAST LINE

The design principle is simple: let the model reason where uncertainty is useful, and move high-cost mistakes into deterministic, observable boundaries.